

Suspicion Scoring Console: A Field Manual for Financial-Crime Analysts

USER MANUAL

AUGUST 2026

everything you need to run a case, end to end



Contents

	How to read this manual	3
1	Getting in	4
2	The Command Center	5
3	Case Registry: creating and managing cases	6
4	The case file: profile and data intake	7
5	The evidence pack: reading the verdict	8
6	Summary drivers: why the score is what it is	9
7	Data quality: how much to trust the ledger	10
8	Rule engine: seventeen typology tests	11
9	Sanctions screening	12
10	AI analyst: narrative with receipts	13
11	Disposition: recording your decision	14
12	Run history: watching the score over time	15
13	Six more views in the evidence pack	16
14	Portfolio analytics: rules and counterparties	17
15	Personalization and the About page	18
16	Field notes: tips and troubleshooting	19
	About ASIA Consulting	20

Disclaimer

This manual documents the ASIA AML Suspicion Scoring Console as deployed for analyst training and case triage. Screens shown use demonstration data; no real customer information appears in this document. Suspicion scores are decision support, not verdicts - the analyst remains responsible for every disposition.

(c) 2026 ASIA Consulting and Private Training. All rights reserved.

How to read this manual

The console reads raw bank statements the way a forensic analyst does, then shows its work. This manual walks every screen in the order you will meet it on a real case.

The ASIA Suspicion Scoring Console turns a stack of bank statement exports into one consolidated ledger, tests that ledger against seventeen money-laundering typology rules, screens names against sanctions lists, and produces a single suspicion probability with a full driver breakdown. An AI assistant then drafts the case narrative - and a verification layer checks every citation it makes against the real ledger before you see it.

Sections 1 to 5 cover the working loop: sign in, watch the portfolio, create a case, feed it data, and read the verdict. Sections 6 to 13 open each evidence view in turn. Sections 14 to 16 cover portfolio analytics, personalization, and field advice.

THE HANDWRITTEN LAYER

Every screenshot in this manual is a real capture of the console. The soft handwritten notes and curved arrows are ours - they point at the exact control being discussed, like a colleague marking up a printout next to you:

blue notes explain what a control does

warm notes flag what deserves caution

BEFORE YOU START


- 1 You need a modern browser and the console address issued by your administrator.
- 2 You need an access code. Codes are issued per team; the console refuses entry without one.
- 3 Have your subject's bank statement exports ready as Excel or CSV files - any of the five supported bank dialects, Arabic column labels included.

1 Getting in

One gate, one code. The landing page states what the system is and asks for exactly one thing.

The console is a restricted system. The landing page carries the ASIA mark, a short statement of what the engine does, and a single access field. Type the code your administrator issued and press ACCESS. There is no self-registration and no password reset flow by design - access is managed centrally.

FIGURE 1 The access gate. Four capability badges under the headline summarise what the engine will do to every case you open.



IF THE GATE DOES NOT OPEN

A rejected code shakes the field and stays on the page. Check for a trailing space, then confirm the code with your administrator. Codes are case-sensitive.

2 The Command Center

Your portfolio at a glance: volumes ingested, risk distribution, and the freshest analysis runs - each one click from its case.

After the gate you land on the Command Center. The left rail lists the five modules of the console; the header carries global search and, on the right, notifications and the settings menu where appearance themes live.

FIGURE 2 | The Command Center. Four KPIs summarise the portfolio; the right panel lists recent scoring runs with their bands.



the five modules - this manual walks them top to bottom

portfolio totals: subjects, files, transactions, gross value

how many cases sit in each risk band

newest runs first - Review jumps straight into the case file

demo helper: loads a sample case to practice on

Dashboard

Portfolio overview and recent runs

Case Registry

Create, open and manage cases

Rule Performance

Which rules fire, portfolio-wide

Counterparty Intel

Parties shared across cases

3 Case Registry: creating and managing cases

One row per subject under investigation. The latest run's band and probability travel with the row.

FIGURE 3 | The Case Registry. Each row shows how many banks, files and transactions a case holds, plus its latest scoring result.

The screenshot shows the ASIA Case Registry interface. On the left is a sidebar with modules: Dashboard, Case Registry (selected), Rule Performance, Counterparty Intel, and About the Application. The main area is titled 'Case Registry' with the subtitle 'Manage subject investigations and view scoring summaries'. It features a search bar and a '+ New Case' button. Below is a table with columns: SUBJECT NAME, REF ID, BANKS, FILES, TXNS, LATEST RUN, and UPDATED. Two rows are visible: 'Ahmad Alsaber' and 'FAHAD'. The 'FAHAD' row has a 'REF ID' of 'DEMO-FAHAD' and a 'LATEST RUN' of 'ELEVATED 46.1%'. Annotations with arrows point to the '+ New Case' button, the 'LATEST RUN' column, the 'REF ID' column, and a trash icon in the 'UPDATED' column.

SUBJECT NAME	REF ID	BANKS	FILES	TXNS	LATEST RUN	UPDATED
Ahmad Alsaber	-	5	5	177	ELEVATED 42.8%	21 Aug 2026
FAHAD	DEMO-FAHAD	5	5	177	ELEVATED 46.1%	22 Aug 2026

start here - New Case asks only for a name and optional reference

band and probability of the newest run

deletes the case and all of its runs - there is no undo

your own file number, if you keep one

CREATE A CASE IN FOUR STEPS

- 1 Press New Case, enter the subject's name and an optional internal reference, and save.
- 2 Open the case and upload the bank statement exports (Excel or CSV). Drag several files at once - each is parsed on arrival.
- 3 Optionally add the subject's financial disclosure PDF through Self Report; the engine extracts declared income and assets from it.
- 4 Press Run Analysis. Scoring is deterministic: the same data always produces the same score.

The parser reads five bank export dialects - compliance exports, core-banking dumps, query extracts, classic statements and signed-amount reports - and understands Arabic column labels. You do not choose the dialect; detection is automatic, file by file.

4 The case file: profile and data intake

Everything the engine knows about the subject, and every file it read - with per-file parsing receipts.

FIGURE 4 Case header. The SCORED chip confirms at least one completed run; Edit Profile holds occupation, declared income and expected countries.

declared economic profile - the baseline every flow is judged against

keep working notes here - they stay with the case

FIGURE 5 Data intake. One card per parsed file: coverage window, masked accounts, rows parsed and an extraction-quality bar.

quality below 70% makes the engine shrink this file's evidence

the receipts: every parsing decision, row by row

THE THREE INTAKE ACTIONS


- 1 Upload Files - add more statement exports at any time; re-run the analysis afterwards.
- 2 Self Report - attach the subject's signed financial disclosure PDF for declared-versus-observed comparison.
- 3 Run Analysis - executes the eight-stage pipeline and stores a new run in the case history.

5 The evidence pack: reading the verdict

One number, honestly derived: the suspicion probability, its band, and the exact movement from prior to posterior.

Below the intake section sits the evidence pack - the analytical heart of the case file. The header shows the suspicion probability produced by Bayesian evidence aggregation, the risk band it falls into, and the analysis metadata. Every element below it exists to let you take this number apart.

FIGURE 6 | The verdict header. The strip under the gauge shows the movement from the 2.0% base-rate prior to this case's posterior.



starts at the 2% base rate - evidence moves it from there

quality of the underlying data - low quality pulls scores toward the prior

the narrative is AI-written, the number never is

LOW
under 5%

MODERATE
5 - 20%

ELEVATED
20 - 50%

HIGH
50 - 80%

CRITICAL
80% +


Twelve tabs line the bottom of the header - Summary Drivers through History. The next eight sections walk the ones you will use on every case; section 13 summarises the rest. Each tab can be opened directly by adding its name to the case address, for example ?tab=rules.

Summary drivers: why the score is what it is

Every contribution to the score, ordered by impact. Green pulls the probability down; red pushes it up.

The drivers list is the fastest way to understand a case. Each row is one piece of evidence with its log-likelihood contribution: the base-rate prior, fired rules, forensic feature signals, and - when data quality is poor - an explicit suppression entry showing how much the engine discounted weak evidence.

FIGURE 7 Bayesian evidence drivers for the demonstration case. The prior appears as its own row, so nothing about the arithmetic is hidden.



the honest starting point: most subjects are not launderers

red bars push suspicion up - the log-LR number is the exact push

the engine deducting for shaky data - suppression is itself listed

HOW TO READ IT


Read top to bottom: rows are ordered by absolute impact, so the three or four biggest bars usually tell the story. Rules from the same family face diminishing returns - a second hit in one family counts half - so the list cannot be inflated by near-duplicate rules.

7 Data quality: how much to trust the ledger

Six checks, each with a stamp and its evidence. Quality below 0.70 shrinks evidence toward the prior.

A score is only as good as the statements behind it. The engine grades the consolidated ledger on six checks - parse coverage, running-balance integrity, period overlap, cross-file duplicates, account coverage gaps and more - and prints the outcome of each with the exact files and rows that earned it.

FIGURE 8 | The data-quality report for the demonstration case: one WARN, one FAIL, and the receipts for both.



the combined quality score used by the engine

FAIL is not cosmetic - balance breaks mean the ledger may be incomplete or altered

per-bank detail lines name the file at fault

The quality score does more than warn. It feeds straight into scoring: below a quality of 0.70 the engine progressively shrinks evidence toward the base rate, and the deduction appears as its own driver row. A clean 46% on solid data and a shaky 46% on broken data are not the same case - this page is how you tell them apart.

8 Rule engine: seventeen typology tests

Each fired rule is a card: what the pattern is, which statement rows triggered it, and the FATF or Egmont source it descends from.

FIGURE 9 | Fired rules for the demonstration case. The mono excerpt inside each card quotes the engine's own finding verbatim.



WORTH KNOWING

Seventeen deterministic rules cover structuring, funnel accounts, dormancy reactivation, velocity bursts, virtual-asset exposure and more. Rules never overlap-count: family diminishing returns halve a second hit from the same family. A rule that did not fire is absent - the engine does not pad the list.

Sanctions screening

Subject and counterparties against OFAC SDN and the UN Consolidated List - with list freshness printed, not presumed.

FIGURE 10 The sanctions view: list freshness at top, the exposure banner, then per-name match cards with alias and identifier detail.

The screenshot displays the Sanctions Screening Console interface. At the top, a section titled 'LIST FRESHNESS AND AUTO RE-SCREEN' shows the status of two lists: 'OFAC SDN' (19,249 entries - 13 min ago) and 'UN CONSOLIDATED' (1,011 entries - 13 min ago). A 'CHECK AND RE-SCREEN NOW' button is visible. Below this, a red banner indicates 'POTENTIAL SANCTIONS EXPOSURE - ANALYST REVIEW REQUIRED' with 1 EXACT and 1 POSSIBLE match. The subject is identified as 'FAHAD' with 'No matches on either list.' Under 'COUNTERPARTY MATCHES (1)', two matches are shown: 'AL-FAISAL, Abdullah Ibrahim' (EXACT match) and 'SADIO AL-AHDAL, Abdallah Fayzal' (POSSIBLE match). Both matches include detailed identifiers such as SIN, DOB, POB, Gender, and National ID.

both lists auto-refresh every six hours - or force it here

EXACT means identical canonical name. POSSIBLE means shared name parts - read the card before reacting

the subject himself: no match on either list

matched via a listed alias - check DOB and documents next

CAUTION

Name similarity is evidence, not identification. Verify identifiers - date of birth, nationality, documents - before acting on any match. Sanctions matches never alter the suspicion score; they are surfaced for your judgment, not added to the arithmetic.

10 AI analyst: narrative with receipts

An assistant drafts the case story - and a verification layer checks every citation it makes against the real ledger.

FIGURE 11 The AI analyst view. The green strip confirms all 197 citations were verified against case records before display.

The screenshot shows the AI analyst view interface. At the top, a green strip indicates 'EVIDENCE VERIFICATION' with '197 of 197 AI citations verified against case records'. Below this, the 'KYC PROFILE CONSISTENCY' section is highlighted with a red border and labeled 'INCONSISTENT'. It contains a narrative about Fahad's government administrative role and lists several features that are irreconcilable with the profile, such as salary deposits from a single government employer and large inward SWIFT credits. A handwritten note points to the '[txn]' references in the text, stating 'every [txn] reference checked against the ledger - broken citations are flagged, not hidden'. Another note points to the 'KYC PROFILE CONSISTENCY' section, stating 'declared profile versus observed behaviour, argued point by point'. The 'TYPOLOGY SYNTHESIS' section is also visible, with a handwritten note pointing to the 'Structuring / Threshold Avoidance' typology, stating 'each typology gets a strength verdict'. A final note points to the 'POSSIBLE BENIGN EXPLANATIONS' section, stating 'the assistant argues both sides - benign readings are listed too'. The interface includes buttons for 'View 2 Core Transactions' and 'View 2 Core Transactions'.

every [txn] reference checked against the ledger - broken citations are flagged, not hidden

declared profile versus observed behaviour, argued point by point

each typology gets a strength verdict

the assistant argues both sides - benign readings are listed too

WHAT THE AI DOES AND NEVER DOES

- + Drafts the narrative, tests KYC consistency, synthesises typologies, and proposes hypotheses for your review.
- Never produces or adjusts the suspicion score. The number exists before the AI runs; deleting the narrative would not change it.

11 Disposition: recording your decision

Three outcomes, an auto-drafted memo, and a hypothesis-by-hypothesis review - the decision stays human.

FIGURE 12 The disposition view. The memo draft is pre-filled from the evidence pack; hypothesis rows record your read of each AI suggestion.

The screenshot shows the 'RECORD DISPOSITION' interface. At the top, it says 'Register your final decision for this scoring pipeline run.' Below this is the 'AUTO-GENERATED CASE MEMO DRAFT' section, which contains a pre-filled memo skeleton: 'ASIA CONSULTING — AML CASE MEMORANDUM', 'CONFIDENTIAL — ANALYST DRAFT FOR REVIEW', 'DATE: [Insert filing date]', 'PREPARED BY: [Analyst name]', 'CASE REFERENCE: [Case number]', and 'SUBJECT AND SCOPE'. A blue arrow points from the text 'a filing-ready memo skeleton, drafted from the case evidence' to this section. Below the memo draft is the 'DECISION' section with three buttons: 'CLOSE - FALSE POSITIVE', 'ADD TO WATCHLIST', and 'ESCALATE TO SAR'. A blue arrow points from the text 'the three dispositions - your call, recorded with your rationale' to these buttons. Below the decision buttons is the 'AI HYPOTHESIS REVIEW (optional - unset rows are recorded as undetermined)' section. It contains a table with six rows of AI hypotheses. Each row has three buttons: 'ACCEPTED', 'DISMISSED', and 'UNDETERMINED'. A blue arrow points from the text 'accept or dismiss each AI hypothesis - unset rows record as undetermined' to the 'UNDETERMINED' buttons. Below the table is the 'RATIONALE & NOTES' section with a text area for 'Detail the reasoning for your decision...'. A blue arrow points from the text 'write why - the rationale is stored with the run' to this section. The bottom of the interface shows the analyst's name 'O. Analyst Level 3' and the company name 'ASIA CONSULTING AND PRIVATE TRAINING - ASIA DATA-SCIENCE'.

RECORD DISPOSITION
Register your final decision for this scoring pipeline run.

AUTO-GENERATED CASE MEMO DRAFT

ASIA CONSULTING — AML CASE MEMORANDUM
CONFIDENTIAL — ANALYST DRAFT FOR REVIEW

DATE: [Insert filing date]
PREPARED BY: [Analyst name]
CASE REFERENCE: [Case number]

SUBJECT AND SCOPE

DECISION

CLOSE - FALSE POSITIVE ADD TO WATCHLIST ESCALATE TO SAR

AI HYPOTHESIS REVIEW (optional - unset rows are recorded as undetermined)

AI-HYP-01 - Unexplained large SWIFT inflows represent proceeds from an undisclosed foreign asset, business...	ACCEPTED	DISMISSED	UNDETERMINED
AI-HYP-02 - Multi-bank funnel structure: many counterparties aggregate funds through Fahad's accounts for...	ACCEPTED	DISMISSED	UNDETERMINED
AI-HYP-03 - Transfer-funded cash extraction at Gulf Bank and Commercial Bank of Kuwait constitutes a mult...	ACCEPTED	DISMISSED	UNDETERMINED
AI-HYP-04 - Dormant account reactivation with a 31,878 KWD credit followed by onward dispatch represents ...	ACCEPTED	DISMISSED	UNDETERMINED
AI-HYP-05 - Virtual-asset exposure alongside large unexplained inflows creates an unobservable layering p...	ACCEPTED	DISMISSED	UNDETERMINED
AI-HYP-06 - The 2021-03-04 split cash deposits represent ITM machine-limit compliance rather than deliber...	ACCEPTED	DISMISSED	UNDETERMINED

RATIONALE & NOTES

Detail the reasoning for your decision...

O. Analyst
Level 3

ASIA CONSULTING AND PRIVATE TRAINING - ASIA DATA-SCIENCE


SHAYNA TOWER, FL 10, RUWAI CITY +966 2227 1724 INFO@ACS-KW.COM

Decisions feed the Rule Performance module: once cases carry dispositions, that view shows which rules concentrate true escalations and which fire mostly on cases that get closed - the evidence base for tuning rule weights over time.

12 Run history: watching the score over time

Every scoring run for the case, on one line. Newer data should move the score for explainable reasons - unexplained drift is itself a finding.

FIGURE 13 Run history for the demonstration case: nineteen runs, one visible score step when new data arrived.



WHEN TO RE-RUN

- 1 After uploading new statement files or a self-report PDF.
- 2 After editing the subject profile - declared income changes the baseline for several rules.
- 3 Periodically on open cases: sanctions lists refresh every six hours, and engine updates are versioned into new runs.

Six more views in the evidence pack

The remaining tabs, and the moment each one earns its place in your review.

VIEW	WHAT IT SHOWS	WHEN TO OPEN IT
Feature Matrix	More than twenty forensic signals - near-threshold density, pass-through ratio, dwell time, Benford deviation - each with its computed value and rating.	When you want the raw measurements behind a rule or a driver row.
Consolidation	How five bank files became one ledger: detected internal transfers, cross-file duplicates and the netting decisions, all reported with row references.	When a total looks off, or you suspect the same money was counted twice.
Timeline	The subject's complete flow history on one axis, flagged windows highlighted.	When you need to see bursts, dormancy and regime shifts at a glance.
Network	A graph of accounts and counterparties around the subject, sized by volume.	When tracing where funnel inflows come from and where value exits.
Investigation Intel	Investigator-oriented leads: repeated references, round-amount chains, counterparties worth a database check.	When preparing the follow-up work that goes beyond the statements.
Transactions	The full consolidated ledger - filter by bank, flagged-only, free text, or a list of transaction ids.	When verifying any citation - rules and the AI link here pre-filtered.


DEEP LINKS

Every tab is addressable. Append ?tab= and the tab name to a case address - rules, sanctions, features, consolidation, timeline, network, intelligence, ai, transactions, disposition or history - to open a case directly on that view. Useful in emails and case-review notes.

14 Portfolio analytics: rules and counterparties

Two modules that look across cases: which rules do the work, and which parties keep reappearing.


FIGURE 14 Rule Performance. Once dispositions are recorded, the split bar shows where each rule's cases ended up.



fires portfolio-wide, latest run per case

fills as your team records decisions - empty is honest, not broken

FIGURE 15 Counterparty Intel. Cards are parties seen in more than one case; the chips underneath jump straight into each case.



same party, two cases - exactly what structuring networks look like

click a chip to open that case

masked references merge only on literal equality - clustering stays conservative

15 Personalization and the About page

Six appearance themes for long shifts and bright rooms - and a methodology dossier that spells out the mathematics.

FIGURE 16 The Paper theme: the same Command Center, re-inked for print-like reading. Appearance never changes data or scores.





FIGURE 17 About the Application: the methodology dossier. Seventeen rules, the scoring mathematics, AI boundaries and honest limitations.



16 Field notes: tips and troubleshooting

The habits that make the console fast - and the first thing to check when something looks wrong.

WORKING HABITS

- 1 Read drivers before narrative. The top three bars usually tell the story; the AI text is the elaboration, not the source.
- 2 Check data quality before trusting a score. A FAIL on running-balance integrity changes how far you lean on the ledger.
- 3 Use View Txns links everywhere. Every rule and citation can be traced to statement rows in two clicks - do the trace for anything you plan to write in a memo.
- 4 Record dispositions promptly. Rule Performance only becomes useful once decisions accumulate.

IF SOMETHING LOOKS WRONG

- Q A file parsed fewer rows than expected. Open View Parser Logs on its intake card - sparse spreadsheet grids legitimately produce few rows, and the log says exactly what was skipped and why.
- Q The score changed between runs. Open History: a quality change, new files, or an engine version explains it. Unexplained drift is worth escalating.
- Q A sanctions match looks impossible. Check the tier - POSSIBLE means shared name parts only. Verify identifiers before reacting; the score was never touched by it.
- Q Totals differ from a bank's own statement summary. Open Consolidation - internal transfers between the subject's accounts are netted, and each netting decision is listed with its rows.

A LAST WORD

The console is built on one promise: no black boxes. If you cannot see why a number is what it is, that is a defect - not a feature to work around. Every screen in this manual exists to keep that promise; use the receipts.

**COMMITTED TO
EVIDENCE-BASED
FINANCIAL-CRIME ANALYTICS**

ASIA Consulting and Private Training is a Kuwait-based consultancy. Its ASIA Data-Science division builds decision-support systems for financial-crime analysis, pairing deterministic forensic engines with carefully bounded AI assistance.

The Suspicion Scoring Console documented here fuses cross-bank statement forensics, typology rules descended from FATF and Egmont guidance, sanctions screening and verified AI narration into one evidence pack that compliance teams can defend.

ASIA Consulting and Private Training

ASIA Data-Science
Shayma Tower, Floor 10
Kuwait City, Kuwait

Tel.: +965 2227 1724
INFO@ACS-KW.COM